

LAB - 9

121CS1133

1) k-Means on a randomly generated dataset

- Keep 6000 random points \
- The number of centers to generate should be initially 4 \
- The standard deviation of the clusters should be 0.9 \

a) Display the scatter plot of the randomly generated data.

b) Initialize the KMeans model and display the plot after k-means clustering.

c) Try to cluster the above random dataset into 3 clusters. and display the plot after clustering.

Ques- What is the optimal number of clusters? Use appropriate method/methods to find optimal number of clusters.

In [10]:

```
import numpy as np
import matplotlib.pyplot as plt

def generate_data(num_points, num_clusters, std_dev):
    np.random.seed(42)
    centers = np.random.rand(num_clusters, 2) * 10
    data = []

    for center in centers:
        points = np.random.normal(loc=center, scale=std_dev, size=(num_points // num_clusters, 2))
        data.append(points)

    return np.vstack(data), centers

def euclidean_distance(a, b):
    return np.sqrt(np.sum((a - b) ** 2))

def k_means(data, k, max_iters=100):
    np.random.seed(42)
    initial_indices = np.random.choice(data.shape[0], k, replace=False)
    centroids = data[initial_indices]

    for _ in range(max_iters):
        labels = np.zeros(data.shape[0])
        for i in range(data.shape[0]):
            distances = [euclidean_distance(data[i], centroid) for centroid in centroids]
            labels[i] = np.argmin(distances)

        new_centroids = np.zeros((k, data.shape[1]))
        for i in range(k):
            points = data[labels == i]
            if len(points) > 0:
                new_centroids[i] = np.mean(points, axis=0)
```

```

    if np.all(centroids == new_centroids):
        break
    centroids = new_centroids

    return labels, centroids

def plot_clusters(data, labels, centroids, title):
    plt.figure(figsize=(10, 6))
    for i in np.unique(labels):
        plt.scatter(data[labels == i][:, 0], data[labels == i][:, 1], s=10)
    plt.scatter(centroids[:, 0], centroids[:, 1], c='red', marker='X', s=200)
    plt.title(title)
    plt.xlabel("X-axis")
    plt.ylabel("Y-axis")
    plt.grid(True)
    plt.show()

```


In [14]:

```


# 1.a)
num_points = 6000
num_initial_clusters = 4
std_dev = 0.9
data, centers = generate_data(num_points, num_initial_clusters, std_dev)

plt.figure(figsize=(10, 6))
plt.scatter(data[:, 0], data[:, 1], s=10)
plt.title("Scatter Plot of Randomly Generated Data")
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.grid(True)
plt.show()


```



```
In [13]: # 1.b)
labels_initial, centroids_initial = k_means(data, num_initial_clusters)
plot_clusters(data, labels_initial, centroids_initial, f'K-Means Clustering with {num_i
```



```
In [15]: #1.c)
labels_three, centroids_three = k_means(data, 3)
plot_clusters(data, labels_three, centroids_three, 'K-Means Clustering with 3 Clusters'
```




In [16]:

```
# Question (Elbow method)
wcss = []
for k in range(1, 11):
    labels, centroids = k_means(data, k)
    wcss.append(sum([euclidean_distance(data[i], centroids[int(labels[i])]) ** 2 for i in range(len(data))]))

plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of Clusters')
plt.ylabel('WCSS (Within-Cluster Sum of Squares)')
plt.grid(True)
plt.show()

optimal_clusters = np.argmin(np.diff(np.diff(wcss))) + 2
print(f"The optimal number of clusters is: {optimal_clusters}")
```



The optimal number of clusters is: 5

In [37]:

```
# k-means++

import numpy as np
import matplotlib.pyplot as plt

def generate_data(num_points, num_clusters, std_dev):
    np.random.seed(42)
    centers = np.random.rand(num_clusters, 2) * 10
    data = []

    for center in centers:
        points = np.random.normal(loc=center, scale=std_dev, size=(num_points // num_clusters, 2))
        data.append(points)

    return np.vstack(data), centers

def euclidean_distance(a, b):
    return np.sqrt(np.sum((a - b) ** 2))

# k-means++ initialization
def initialize_centroids(X, k):
    centroids = np.zeros((k, X.shape[1]))
    centroids[0] = X[np.random.choice(X.shape[0])]

    for i in range(1, k):
        distances = np.array([min([euclidean_distance(x, c) for c in centroids[:i]]) for x in X])
        probabilities = distances / distances.sum()
        centroids[i] = X[np.random.choice(X.shape[0], p=probabilities)]

    return centroids

def k_means(X, k, max_iters=100):
    centroids = initialize_centroids(X, k)
```

```

for _ in range(max_iters):
    labels = np.zeros(X.shape[0])
    for i in range(X.shape[0]):
        distances = [euclidean_distance(X[i], centroid) for centroid in centroids]
        labels[i] = np.argmin(distances)

    new_centroids = np.zeros((k, X.shape[1]))
    for i in range(k):
        points = X[labels == i]
        if len(points) > 0:
            new_centroids[i] = np.mean(points, axis=0)

    if np.all(centroids == new_centroids):
        break
    centroids = new_centroids

return labels, centroids

def plot_clusters(X, labels, centroids, title):
    plt.figure(figsize=(10, 6))
    for i in np.unique(labels):
        plt.scatter(X[labels == i][:, 0], X[labels == i][:, 1], s=10, label=f'Cluster {i}')
    plt.scatter(centroids[:, 0], centroids[:, 1], c='red', marker='X', s=200, label='Centroids')
    plt.title(title)
    plt.xlabel("X-axis")
    plt.ylabel("Y-axis")
    plt.legend()
    plt.grid(True)
    plt.show()

# a)
num_points = 6000
num_initial_clusters = 4
std_dev = 0.9
data, centers = generate_data(num_points, num_initial_clusters, std_dev)

# Display the scatter plot
plt.figure(figsize=(10, 6))
plt.scatter(data[:, 0], data[:, 1], s=10)
plt.title("Scatter Plot of Randomly Generated Data")
plt.xlabel("X-axis")
plt.ylabel("Y-axis")
plt.grid(True)
plt.show()

# b)
labels_initial, centroids_initial = k_means(data, num_initial_clusters)
plot_clusters(data, labels_initial, centroids_initial, f'K-Means Clustering with {num_initial_clusters} Clusters')


# c)
labels_three, centroids_three = k_means(data, 3)
plot_clusters(data, labels_three, centroids_three, 'K-Means Clustering with 3 Clusters')



# Determine the optimal number of clusters using the elbow method
wcss = []
for k in range(1, 11):
    labels, centroids = k_means(data, k)
    wcss.append(sum([euclidean_distance(data[i], centroids[int(labels[i])]) ** 2 for i in range(data.shape[0])]))


```

```
plt.figure(figsize=(10, 6))
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method for Optimal k')
plt.xlabel('Number of Clusters')
plt.ylabel('WCSS (Within-Cluster Sum of Squares)')
plt.grid(True)
plt.show()

# Optimal number of clusters based on the elbow method
optimal_clusters = np.argmin(np.diff(np.diff(wcss))) + 2
print(f"The optimal number of clusters is: {optimal_clusters}")
```







The optimal number of clusters is: 9

2) k-Means on country-continent dataset

- Read the dataset from the CSV file.
- Get the unique continents from the dataset.
- Map text data to numbers.
- Run the k-means algorithm with the number of continents clusters.
- Plot the results.

Ques- Compare the results with different number of clusters. What do you observe?

```
In [31]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [ ]: # a) Read the dataset
file_path = 'countries_continents.csv'
data = pd.read_csv(file_path)
```

```
In [ ]: # b) Get unique continents
unique_continents = data['Continent'].unique()
print("Unique Continents:", unique_continents)
```

Unique Continents: ['North America' 'Asia' 'Africa' 'Europe' 'South America' 'Oceania' 'Antarctica' 'Seven seas (open ocean)']

```
In [ ]: # c) Map text to numbers
continent_mapping = {continent: idx for idx, continent in enumerate(unique_continents)}
print("Continent Mapping:", continent_mapping)

continent_numbers = []
for continent in data['Continent']:
    continent_numbers.append(continent_mapping[continent])

data['Continent_Num'] = continent_numbers
print(data.head())
```

Continent Mapping: {'North America': 0, 'Asia': 1, 'Africa': 2, 'Europe': 3, 'South America': 4, 'Oceania': 5, 'Antarctica': 6, 'Seven seas (open ocean)': 7}

	Country	Longitude	Latitude	Continent	Continent_Num
0	Aruba	-69.982677	12.520880	North America	0
1	Afghanistan	66.004734	33.835231	Asia	1
2	Angola	17.537368	-12.293361	Africa	2
3	Anguilla	-63.064989	18.223959	North America	0
4	Albania	20.049834	41.142450	Europe	3

```
In [ ]: # d) k-means algorithm
X = data[['Longitude', 'Latitude']].values

def k_means(X, k, max_iters=100):
    np.random.seed(42)
    centroids = X[np.random.choice(X.shape[0], k, replace=False)]

    for _ in range(max_iters):
        distances = np.linalg.norm(X[:, np.newaxis] - centroids, axis=2)
        labels = np.argmin(distances, axis=1)



        new_centroids = np.array([X[labels == i].mean(axis=0) for i in range(k)])



        if np.all(centroids == new_centroids):
            break
        centroids = new_centroids


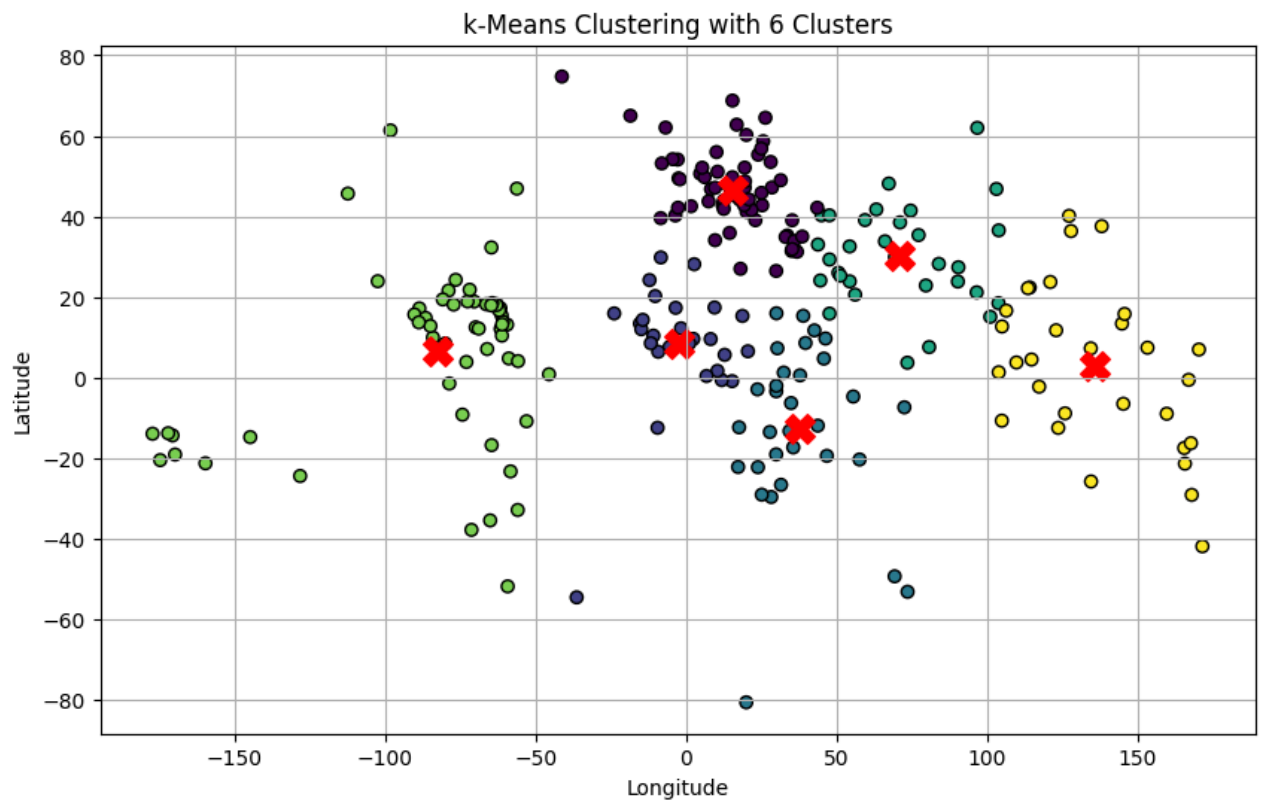
    return labels, centroids
```


```
In [35]: # e) Plotting
def plot_clusters(X, labels, centroids, title):
    plt.figure(figsize=(10, 6))
    plt.scatter(X[:, 0], X[:, 1], c=labels, cmap='viridis', marker='o', edgecolor='k')
    plt.scatter(centroids[:, 0], centroids[:, 1], c='red', marker='X', s=200)
    plt.title(title)
    plt.xlabel('Longitude')
    plt.ylabel('Latitude')
    plt.grid(True)
    plt.show()

for k in range(2, len(unique_continents) + 1):
    labels, centroids = k_means(X, k)
    plot_clusters(X, labels, centroids, f'k-Means Clustering with {k} Clusters')
```









e)

Observations:

1. With fewer clusters ($k=2$ to $k=3$), continents are grouped closer together.
2. As the number of clusters increases ($k=4$ to $k=6$), distinct continents are better separated.
3. At higher cluster counts ($k=7$ to $k=8$), clusters align closely with actual continents.
4. The 'Seven seas (open ocean)' is treated as a distinct cluster, indicating its geographical separation.